Flow Elasticity in ETFs:
A Distributed-Lag Framework for Measuring How Investor Flows Respond to Prior Returns Across Exposure Types

NASDAQ ETF Intelligence / Indexes
Internal Project Write-Up by Clete Roberts
Idea Proposal by Ryan Green

March 13, 2026

Abstract

This project proposes a mathematically rigorous framework for estimating *flow elasticity* in exchange-traded funds (ETFs): the sensitivity and timing of investor flows with respect to prior fund performance. The central hypothesis is that the speed and magnitude of flow response differ systematically by exposure type. For example, commodity-linked exposures such as gold may exhibit faster and more state-dependent flow reactions than broad equity market exposures, reflecting distinct investor motives, including inflation hedging, risk aversion, tactical allocation, and momentum-chasing behavior.

We develop a panel-data and distributed-lag methodology to model future flows as a function of lagged returns, relative performance, and fund characteristics. The project also introduces a tractable definition of a *momentum score* based on estimated dynamic flow response coefficients. The resulting framework can be used for ETF intelligence, issuer analytics, category monitoring, product design, and index research.

Contents

1 Project Motivation **4**

2 Core Research Questions **4**

3 Conceptual Economic Hypothesis **4**
3.1 Broad Market ETFs 5
3.2 Gold ETFs 5
3.3 Thematic and High-Beta Exposures 5
3.4 Fixed Income ETFs 5

4 Project Objective **5**

5 Data Requirements **5**
5.1 Core ETF-Level Data 6
5.2 Derived Variables 6
5.3 External / Macro Data That Could Be Useful 7

6 Flow Construction and Measurement Issues **7**
6.1 Dollar Flows vs. AUM Changes . . . 7
6.2 Normalized Flows . . . 7
6.3 Winsorization / Tail Control . . . 8

7 Initial Exploratory Framework **8**
7.1 Bucketed Return-Flow Analysis . . . 8
7.2 Block Return Idea from the Conversation . . . 8

8 Baseline Econometric Specification **9**
8.1 Interpretation . . . 9
8.2 Why Fixed Effects Matter . . . 9

9 Distributed-Lag Regression **9**
9.1 Model . . . 9
9.2 Economic Meaning . . . 10
9.3 Cumulative Response . . . 10

10 Exposure-Specific Estimation **10**
10.1 Separate Regressions by Exposure . . . 10
10.2 Pooled Regression with Interactions . . . 10

11 Relative Performance vs. Absolute Performance **11**
11.1 Why Relative Performance May Matter . . . 11
11.2 Relative Return Variable . . . 11
11.3 Recommendation . . . 11

12 A Mathematical Definition of Flow Elasticity **11**
12.1 Interpretation . . . 11
12.2 Elasticity as a Momentum Score . . . 12

13 Nonlinearity and Threshold Effects **12**
13.1 Convexity . . . 12
13.2 Piecewise or Spline Models . . . 12
13.3 Asymmetry . . . 13

14 Econometric Challenges and Mitigations **13**
14.1 Persistence and Serial Correlation . . . 13
14.2 Endogeneity . . . 13
14.3 Small Fund Distortion . . . 13
14.4 Survivorship Bias . . . 13
14.5 Exposure Misclassification . . . 13

15 Empirical Evaluation and Model Comparison **13**
15.1 Model Ablation Framework . . . 14
15.2 Cross-Exposure Comparison . . . 14
15.3 Predictive Performance Evaluation . . . 15
15.4 Interpretation of Results . . . 16

16 Suggested Model Sequence **16**

17 Example Formal Specifications **17**
17.1 Specification A: Simple Trailing Return Model . . . 17
17.2 Specification B: Recent vs. Older Block Returns . . . 17
17.3 Specification C: Distributed-Lag Model . . . 17
17.4 Specification D: Exposure Interaction Model . . . 17
17.5 Specification E: Nonlinear Flow Elasticity . . . 17

18 Interpretation of Possible Outcomes **17**
18.1 Case 1: Gold Has Faster Response . . . 17
18.2 Case 2: Broad Market Has Slower Cumulative Response . . . 17
18.3 Case 3: Thematic ETFs Show Convexity . . . 18
18.4 Case 4: Relative Performance Dominates . . . 18

19 Potential Deliverables for NASDAQ ETF Intelligence / Indexes **18**
19.1 Internal Research Note . . . 18
19.2 Exposure-Level Flow Elasticity Dashboard . . . 18
19.3 Fund-Level Predicted Flow Signal . . . 18
19.4 Index/Product Intelligence . . . 18

20 Data Table Blueprint **18**

21 Recommended Empirical Design Choices **20**
21.1 Frequency . . . 20
21.2 Sample Horizon . . . 20
21.3 Category Construction . . . 20
21.4 Inference . . . 20

22 Why This Project Matters Strategically **20**

23 Final Recommendation **21**

24 Very Simple Project Roadmap **21**
24.1 Step 1: Build the ETF dataset . . . 21
24.2 Step 2: Create the main project variables . . . 22
24.3 Step 3: Do simple descriptive analysis first . . . 22
24.4 Step 4: Run the basic regression . . . 22
24.5 Step 5: Run the distributed-lag regression . . . 23
24.6 Step 6: Compare categories and exposure types . . . 23
24.7 Step 7: Add relative performance and nonlinear effects . . . 23
24.8 Step 8: Build the flow elasticity / momentum score . . . 24
24.9 Step 9: Turn the model into an internal research product . . . 24
24.10 Simple Summary of the Order . . . 24

1 Project Motivation

A recurring empirical question in ETF markets is not merely whether flows follow returns, but *how quickly*, *how strongly*, and *under what exposure-specific conditions* they do so. A broad-market equity ETF and a gold ETF may both experience inflows after strong trailing returns, yet the mechanism governing those inflows may differ sharply.

Broad equity exposures may receive more gradual allocations associated with strategic rebalancing, retirement contributions, model portfolio implementation, and delayed advisor adoption. By contrast, gold exposures may reflect more immediate tactical reactions to inflation concerns, macro uncertainty, or crisis hedging demand. Thus, the appropriate object of study is a dynamic response function:

$$\text{future flows} = f(\text{lagged returns}, \text{exposure type}, \text{market regime}, \text{fund characteristics}).$$

The goal is to quantify this response in a manner that is:

(i) statistically credible,

(ii) economically interpretable,

(iii) useful for ETF intelligence work,

(iv) comparable across categories and exposures,

(v) extensible into a ranking or "momentum score."

2 Core Research Questions

This project is designed to answer the following questions:

Q1. Do ETF flows respond positively to prior returns?

Q2. How does the *speed of adjustment* differ across exposure types (e.g., gold, broad U.S. equity, sector, thematic, international equity, fixed income)?

Q3. Is the relationship between flows and prior returns linear, or does it become stronger in extreme return environments?

Q4. Does *relative performance within category* matter more than absolute performance?

Q5. Can we estimate a robust *flow elasticity* measure that can serve as an internal ETF momentum or investor-demand score?

Q6. Are some exposures characterized by short-horizon tactical demand, while others exhibit slower medium-horizon accumulation?

3 Conceptual Economic Hypothesis

The working intuition is that different ETF categories encode different investor use-cases.

3.1 Broad Market ETFs

Broad market ETFs are often used for strategic asset allocation, long-term beta exposure, and model portfolio implementation. Flows may respond to prior returns, but the response is often slower and more persistent because allocation decisions may pass through advisors, investment committees, or periodic rebalancing schedules.

3.2 Gold ETFs

Gold ETFs may respond more rapidly to prior performance because investor demand is often linked to macro narratives such as inflation, real-rate moves, crisis hedging, or defensive positioning. This suggests a stronger near-term elasticity:

$$\frac{\partial \text{flows}_{t+h}}{\partial \text{return}_{t-k}}$$

may be concentrated at shorter lags for gold than for broad market ETFs.

3.3 Thematic and High-Beta Exposures

Thematic, leveraged, or speculative exposures may exhibit highly convex flow responses. Strong recent returns can attract rapid retail attention, while poor returns may induce equally sharp redemptions. The response may be asymmetric and nonlinear.

3.4 Fixed Income ETFs

Fixed income flows may depend less on raw trailing returns and more on yield levels, duration risk, credit spreads, and macro policy expectations. For these exposures, the return-flow relationship may exist but should likely be augmented with rate and spread variables.

4 Project Objective

The primary project objective is to estimate, compare, and interpret **ETF flow elasticity** across exposure types.

A useful operational definition is:

> *Flow elasticity* is the expected change in normalized ETF flows resulting from a given change in lagged fund returns, conditional on exposure type and fund characteristics.

We want both:

(a) a **static elasticity estimate** (how sensitive flows are), and

(b) a **dynamic elasticity profile** (how quickly that sensitivity materializes over subsequent months).

5 Data Requirements

A strong version of this project requires a fund-level panel dataset over approximately the past 10 years at monthly frequency, though weekly data could be explored later.

5.1 Core ETF-Level Data

For each ETF i and month t, we ideally require:

D1. Net flows: dollar net creations/redemptions.

D2. AUM: assets under management at month-end.

D3. Returns: total return over monthly horizon.

D4. Price / NAV information: to separate flow-driven AUM changes from market-driven AUM changes.

D5. Category / exposure classification: e.g., gold, broad U.S. equity, sector equity, international equity, fixed income, commodity, leveraged/inverse, thematic.

D6. Fund age: time since launch.

D7. Expense ratio.

D8. Liquidity measures: ADV, bid-ask spread, median daily volume, turnover.

D9. Size measures: AUM levels, log AUM.

D10. Issuer.

D11. Benchmark / index family.

D12. Structural features: physically backed, synthetic, active/passive, single-commodity, broad commodity, sector, etc.

5.2 Derived Variables

From the raw data, we would construct:

V1. Normalized flows:

$$y_{i,t} = \frac{\text{Flow}_{i,t}}{\text{AUM}_{i,t-1}}$$

or alternatively

$$y_{i,t} = \frac{\text{Flow}_{i,t}}{\frac{1}{2}(\text{AUM}_{i,t-1} + \text{AUM}_{i,t})}$$

for scale comparability across funds.

V2. Trailing return windows: 1M, 3M, 6M, 12M returns.

V3. Rolling subperiod return blocks: for example,

$$R_{i,t}^{(0,3)} = \text{return over months } t-2 \text{ to } t,$$

$$R_{i,t}^{(3,6)} = \text{return over months } t-5 \text{ to } t-3.$$

V4. Category-relative performance:

$$\widetilde{R}_{i,t}^{(k)} = R_{i,t}^{(k)} - \text{median return of category } c(i) \text{ over the same window.}$$

V5. Return ranks / percentiles within category.

V6. Market share of category AUM.

V7. Flow share within category.

V8. Lagged volatility, downside volatility, drawdown.

V9. Interaction terms for macro regimes, such as inflationary periods, equity drawdown periods, or high-volatility states.

5.3 External / Macro Data That Could Be Useful

For a stronger economic interpretation, the following may be useful:

M1. Treasury yields and yield changes.

M2. Real yields.

M3. CPI / inflation surprise proxies.

M4. Equity market volatility (e.g., VIX).

M5. Credit spreads.

M6. Gold spot returns / commodity index returns.

M7. Dollar index.

M8. Broad risk-on / risk-off indicators.

These are especially helpful if the project expands beyond a pure return-flow response study into a conditional demand model.

6 Flow Construction and Measurement Issues

A key practical challenge is to ensure that flows are measured correctly.

6.1 Dollar Flows vs. AUM Changes

AUM changes conflate market performance and net creations/redemptions. Therefore, one should not use raw changes in AUM as a proxy for flows unless flows are unavailable and must be inferred.

A standard approximation is:

$$\text{Flow}_{i,t} \approx \text{AUM}_{i,t} - \text{AUM}_{i,t-1}(1 + r_{i,t}),$$

where $r_{i,t}$ is the total return over period t.

6.2 Normalized Flows

Since a \$500 million inflow means very different things for a \$1 billion ETF versus a \$100 billion ETF, normalization is essential:

$$\text{FlowRate}_{i,t} = \frac{\text{Flow}_{i,t}}{\text{AUM}_{i,t-1}}.$$

This also improves comparability across categories and makes regression coefficients easier to interpret.

6.3 Winsorization / Tail Control

Flow rates can be extreme for very small funds. It may therefore be necessary to winsorize the dependent variable, for example at the 1st and 99th percentiles, or exclude observations with extremely small lagged AUM.

7 Initial Exploratory Framework

Before formal modeling, the first stage should be descriptive.

7.1 Bucketed Return-Flow Analysis

A simple and intuitive first pass is:

(1) compute trailing 6M returns for each fund-month,

(2) bucket those returns into deciles or bins,

(3) compute average next-month flow rate by bucket,

(4) repeat by exposure type.

This gives a nonparametric view of:

$$\mathbb{E}[y_{i,t+1} \mid R_{i,t}^{6M} \in \text{bucket } b, \text{ exposure } = e].$$

Such charts will likely reveal whether flow-response is approximately linear, convex, threshold-based, or asymmetric.

7.2 Block Return Idea from the Conversation

A particularly useful decomposition is:

- use months 0–3 performance as a recent return signal,
- use months 3–6 performance as an earlier return signal,
- examine how flows in the next month or next quarter respond differently to those blocks.

Formally:

$$R_{i,t}^{\text{recent}} = \sum_{j=0}^{2} r_{i,t-j}, \qquad R_{i,t}^{\text{older}} = \sum_{j=3}^{5} r_{i,t-j}.$$

Then estimate:

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{\text{recent}} + \beta_2 R_{i,t}^{\text{older}} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1}.$$

This directly tests whether investors respond more to very recent returns than to older but still recent performance.

8 Baseline Econometric Specification

A natural baseline panel regression is:

$$y_{i,t+1} = \alpha_i + \tau_t + \beta R_{i,t}^{(6M)} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1},$$

where:

- $y_{i,t+1}$ is next-month normalized flow,
- α_i are fund fixed effects,
- τ_t are time fixed effects,
- $R_{i,t}^{(6M)}$ is trailing 6-month return,
- $\boldsymbol{X}_{i,t}$ includes controls such as log AUM, age, expense ratio, volatility, and liquidity.

8.1 Interpretation

If $\beta > 0$, then better trailing returns predict larger future inflows on average. This is the first empirical representation of flow elasticity.

8.2 Why Fixed Effects Matter

Fund fixed effects absorb persistent unobserved differences such as brand, issuer strength, benchmark popularity, and structural features. Time fixed effects absorb market-wide shocks affecting all funds simultaneously.

9 Distributed-Lag Regression

The more conceptually powerful model is a distributed-lag specification. This estimates not merely whether returns matter, but *when* they matter.

9.1 Model

Let $r_{i,t-k}$ denote fund return k months ago. Then estimate:

$$y_{i,t} = \alpha_i + \tau_t + \sum_{k=1}^{K} \beta_k r_{i,t-k} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t-1} + \varepsilon_{i,t}.$$

Here:

- β_1 measures the effect of last month's return on current flow,
- β_2 measures the effect of the return from two months ago,
- and so forth up to lag K.

9.2 Economic Meaning

The sequence $\{\beta_k\}_{k=1}^K$ is the **flow response curve**. It provides a dynamic map of investor reaction speed.

If gold ETFs react quickly, we may observe:

$$\beta_1, \beta_2 \text{ large and positive,}$$

with coefficients fading rapidly after that. If broad-market ETFs react more slowly, the mass of coefficients may be more distributed across lags.

9.3 Cumulative Response

Define the cumulative response through lag h as:

$$C(h) = \sum_{k=1}^{h} \beta_k.$$

Then:

- $C(1)$ measures immediate elasticity,
- $C(3)$ measures short-horizon elasticity,
- $C(6)$ measures medium-horizon elasticity.

This is one of the cleanest ways to compare investor response speed across categories.

10 Exposure-Specific Estimation

There are two main ways to compare categories.

10.1 Separate Regressions by Exposure

For each exposure class e, estimate:

$$y_{i,t} = \alpha_i + \tau_t + \sum_{k=1}^{K} \beta_{e,k} r_{i,t-k} + \boldsymbol{\gamma}_e' \boldsymbol{X}_{i,t-1} + \varepsilon_{i,t}.$$

This yields an exposure-specific response profile:

$$\{\beta_{e,1}, \beta_{e,2}, \ldots, \beta_{e,K}\}.$$

10.2 Pooled Regression with Interactions

Alternatively, estimate one pooled model:

$$y_{i,t} = \alpha_i + \tau_t + \sum_{k=1}^{K} \beta_k r_{i,t-k} + \sum_{e \in \mathcal{E}} \sum_{k=1}^{K} \delta_{e,k} \mathbf{1}\{i \in e\} r_{i,t-k} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t-1} + \varepsilon_{i,t}.$$

The interaction terms $\delta_{e,k}$ indicate how the flow-response curve for exposure e differs from the baseline category.

11 Relative Performance vs. Absolute Performance

The conversation correctly notes that outperformance relative to peers may matter, though it adds complexity.

11.1 Why Relative Performance May Matter

Investors do not allocate in a vacuum. Many compare funds within a category. A broad equity ETF that outperformed its category peers may attract inflows even if absolute returns were modest. Conversely, a fund with high absolute returns may still lag peers and fail to attract attention.

11.2 Relative Return Variable

Define:

$$\widetilde{R}_{i,t}^{(6M)} = R_{i,t}^{(6M)} - \overline{R}_{c(i),t}^{(6M)},$$

where $\overline{R}_{c(i),t}^{(6M)}$ is the category average or median.

Then estimate:

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{(6M)} + \beta_2 \widetilde{R}_{i,t}^{(6M)} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1}.$$

11.3 Recommendation

For project phase 1, begin with absolute returns only. Then add relative returns as a phase 2 extension once the core dynamic is established.

12 A Mathematical Definition of Flow Elasticity

A clean internal measure is:

$$\mathcal{E}_e(h) = \sum_{k=1}^{h} \beta_{e,k},$$

where $\beta_{e,k}$ is the estimated effect of a one-unit increase in lagged return at lag k on normalized flows for exposure e.

12.1 Interpretation

If returns are expressed in decimal units, then $\mathcal{E}_e(h)$ can be interpreted as:

> the cumulative change in flow rate over the next h months associated with a 100 percentage point change in return.

For practical interpretation, it is often better to scale by a 1% return:

$$\mathcal{E}_e^{(1\%)}(h) = 0.01 \cdot \mathcal{E}_e(h).$$

Then one can say:

> A 1% higher trailing monthly return is associated with approximately x basis points of cumulative flow-rate response over the next h months.

12.2 Elasticity as a Momentum Score

One can define an exposure-level or fund-level momentum score as:

$$\text{MomentumScore}_e = w_1\mathcal{E}_e(1) + w_2\mathcal{E}_e(3) + w_3\mathcal{E}_e(6),$$

for chosen weights $w_1, w_2, w_3 > 0$.

A faster-reacting category would receive a higher score if weights emphasize near-term response.

An alternative score for a specific fund i at time t is:

$$\text{PredictedFlowScore}_{i,t+1} = \sum_{k=1}^{K} \widehat{\beta}_{c(i),k} r_{i,t-k},$$

which gives the model-implied next-period flow signal based on recent return history and category-specific response coefficients.

13 Nonlinearity and Threshold Effects

It is unlikely that the return–flow relationship is perfectly linear. Investors may respond more strongly to extreme performance, particularly in exposures that attract attention-driven or tactical allocation behavior.

13.1 Convexity

Strong positive returns may trigger disproportionate media attention and retail participation, especially in thematic, single-asset, or commodity exposures. This suggests a convex relationship between past returns and future flows.

We test this using a quadratic specification:

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{(6M)} + \beta_2 (R_{i,t}^{(6M)})^2 + \gamma' X_{i,t-1} + \varepsilon_{i,t+1}.$$

If $\beta_2 > 0$, the marginal sensitivity of flows to past returns increases with return magnitude. The derivative

$$\frac{\partial y}{\partial R} = \beta_1 + 2\beta_2 R$$

shows that larger positive returns generate disproportionately larger flow responses.

13.2 Piecewise or Spline Models

A more flexible alternative is to estimate the relationship using return bins or spline functions:

$$y_{i,t+1} = \alpha_i + \tau_t + \sum_{m=1}^{M} \beta_m B_m(R_{i,t}^{(6M)}) + \gamma' X_{i,t-1} + \varepsilon_{i,t+1},$$

where $B_m(\cdot)$ are spline basis functions.

This approach allows the return–flow relationship to be estimated without imposing a specific functional form such as quadratic convexity and can reveal threshold regions where flows increase sharply after strong performance.

13.3 Asymmetry

Investor responses may also differ for positive and negative returns. For example, investors may chase strong performers more aggressively than they redeem underperforming funds.

This can be tested by separating positive and negative return components:

$$R_{i,t}^{+} = \max(R_{i,t}, 0), \qquad R_{i,t}^{-} = \min(R_{i,t}, 0),$$

and estimating separate coefficients for each component.

14 Econometric Challenges and Mitigations

14.1 Persistence and Serial Correlation

Returns and flows may both be serially correlated. Standard errors should therefore be clustered, ideally at the fund level, and potentially two-way clustered by fund and time.

14.2 Endogeneity

The project is primarily predictive/descriptive rather than strictly causal. A return-flow relation does not necessarily imply that past returns cause flows in a structural sense. However, using lagged returns alleviates contemporaneous simultaneity.

14.3 Small Fund Distortion

Small ETFs can exhibit enormous percentage flows. AUM thresholds, winsorization, or weighted regressions may be needed.

14.4 Survivorship Bias

The ETF universe should include dead, merged, or delisted funds where possible. Otherwise, the sample may overstate the persistence and attractiveness of successful exposures.

14.5 Exposure Misclassification

Category definitions matter. Broad categories should be internally coherent. For example, "gold" should likely distinguish physically backed gold funds from leveraged gold miner equity funds.

15 Empirical Evaluation and Model Comparison

This section outlines the empirical evaluation framework used to assess the explanatory power and robustness of the proposed flow-response model. Because the model contains several potential components — including lagged returns, distributed lags, relative performance, and nonlinear response terms — it is important to evaluate which components meaningfully contribute to explaining ETF flow dynamics.

To address this, we adopt a structured **model ablation and comparison approach**. This involves estimating a sequence of progressively richer model specifications and comparing their explanatory power and response characteristics.

15.1 Model Ablation Framework

Ablation analysis evaluates the marginal contribution of each modeling component by estimating nested specifications that incrementally introduce additional structure.

The sequence of models considered is as follows:

1. **Baseline Model**: Controls only, including fund fixed effects and time fixed effects.
2. **Return Model**: Adds trailing return measures.
3. **Relative Performance Model**: Adds category-relative return measures.
4. **Distributed Lag Model**: Allows flows to respond dynamically to returns over multiple past periods.
5. **Full Model**: Adds nonlinear response terms and additional robustness controls.

This structure allows us to determine which elements of the model meaningfully improve explanatory power and which provide limited incremental value.

Table 1: Planned Model Ablation Structure

Model	Lagged Returns	Relative Returns	Distributed Lag	Nonlinearity	Controls
Baseline	×	×	×	×	✓
Return Model	✓	×	×	×	✓
Relative Model	✓	✓	×	×	✓
Distributed Lag Model	✓	✓	✓	×	✓
Full Specification	✓	✓	✓	✓	✓

For each specification we will report measures such as:

- within-sample R^2
- adjusted R^2
- root mean squared prediction error (RMSE)
- coefficient stability across subsamples

This comparison allows us to evaluate the relative explanatory value of each modeling component.

15.2 Cross-Exposure Comparison

An important objective of this project is to compare how investor flows respond to past returns across ETF exposure types. To facilitate interpretation, the estimated flow-response coefficients will be summarized using cumulative response measures across several horizons.

Specifically, we will compute cumulative response measures of the form

$$\mathcal{R}_e(h) = \sum_{k=1}^{h} \beta_{e,k},$$

which capture the total sensitivity of flows to past returns over the first h lags.

These measures will be compared across major ETF exposure categories, such as:

- gold and commodity exposures
- broad equity market ETFs
- sector ETFs
- thematic exposures
- fixed income ETFs

Table 2 illustrates the planned comparison structure.

Table 2: Illustrative Cross-Exposure Flow Response Comparison

Exposure Type	Immediate Response	3-Month Response	6-Month Response	Peak Lag
Gold ETFs	–	–	–	–
Broad Equity ETFs	–	–	–	–
Sector ETFs	–	–	–	–
Thematic ETFs	–	–	–	–
Fixed Income ETFs	–	–	–	–

These comparisons allow us to identify differences in both the magnitude and speed of investor flow reactions across exposures.

15.3 Predictive Performance Evaluation

In addition to coefficient interpretation, we will evaluate the predictive performance of the proposed model using several forecast accuracy metrics.

Specifically, we will compare model variants using:

- root mean squared error (RMSE)
- mean absolute error (MAE)
- directional accuracy of predicted flow changes
- out-of-sample predictive R^2

Table 3 illustrates the planned comparison structure.

Table 3: Illustrative Predictive Performance Comparison

Model	RMSE	MAE	R^2	Direction Accuracy
Baseline Model	–	–	–	–
Return Model	–	–	–	–
Distributed Lag Model	–	–	–	–
Full Specification	–	–	–	–

This evaluation framework ensures that the proposed specification is justified not only by conceptual motivation but also by measurable improvements in empirical performance.

15.4 Interpretation of Results

The combination of ablation analysis, exposure-level comparisons, and predictive performance metrics provides a comprehensive evaluation of the proposed flow-response framework.

This approach allows us to answer three central empirical questions:

1. Which model components meaningfully improve the explanatory power of ETF flow models?
2. How does the magnitude and timing of flow responses differ across exposure types?
3. To what extent can historical return dynamics help predict near-term ETF flow pressure?

By structuring the analysis in this way, the results can be interpreted both as a behavioral description of investor allocation dynamics and as a practical forecasting framework for ETF demand.

16 Suggested Model Sequence

A practical project roadmap would proceed in stages.

Stage 1: Descriptive Analytics

(1) Construct monthly normalized fund flows.

(2) Group ETFs by exposure type.

(3) Plot next-month average flows versus trailing 6M return buckets.

(4) Compare shapes across exposures.

Stage 2: Baseline Regressions

(1) Estimate fixed-effects panel models with trailing 1M, 3M, and 6M returns.

(2) Add controls: size, age, expense ratio, liquidity, volatility.

(3) Evaluate coefficient stability across subperiods.

Stage 3: Distributed-Lag Estimation

(1) Estimate lag structure over $K = 6$ or $K = 12$ months.

(2) Compute cumulative response curves.

(3) Compare speed of adjustment across exposure classes.

Stage 4: Nonlinearity and Relative Performance

(1) Add relative performance measures within category.

(2) Test nonlinear and asymmetric response.

(3) Build predicted-flow or momentum scores.

Stage 5: Production / Intelligence Output

(1) Create exposure-level rankings by flow elasticity.

(2) Generate fund-level expected-flow dashboards.

(3) Integrate insights into ETF intelligence commentary or index/product research.

17 Example Formal Specifications

17.1 Specification A: Simple Trailing Return Model

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{1M} + \beta_2 R_{i,t}^{3M} + \beta_3 R_{i,t}^{6M} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1}.$$

17.2 Specification B: Recent vs. Older Block Returns

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{(0,3)} + \beta_2 R_{i,t}^{(3,6)} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1}.$$

This directly captures whether investors react more strongly to the most recent 3 months than to the prior 3 months.

17.3 Specification C: Distributed-Lag Model

$$y_{i,t} = \alpha_i + \tau_t + \sum_{k=1}^{6} \beta_k r_{i,t-k} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t-1} + \varepsilon_{i,t}.$$

17.4 Specification D: Exposure Interaction Model

$$y_{i,t} = \alpha_i + \tau_t + \sum_{k=1}^{6} \beta_k r_{i,t-k} + \sum_{e \in \mathcal{E}} \sum_{k=1}^{6} \delta_{e,k} \mathbf{1}\{i \in e\} r_{i,t-k} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t-1} + \varepsilon_{i,t}.$$

17.5 Specification E: Nonlinear Flow Elasticity

$$y_{i,t+1} = \alpha_i + \tau_t + \beta_1 R_{i,t}^{(6M)} + \beta_2 \left(R_{i,t}^{(6M)} \right)^2 + \beta_3 \widetilde{R}_{i,t}^{(6M)} + \boldsymbol{\gamma}' \boldsymbol{X}_{i,t} + \varepsilon_{i,t+1}.$$

18 Interpretation of Possible Outcomes

18.1 Case 1: Gold Has Faster Response

Suppose gold ETFs exhibit a large positive β_1 and β_2, with little remaining mass after lag 2. This would suggest investors react quickly to recent performance in gold, consistent with tactical demand or macro-sensitive positioning.

18.2 Case 2: Broad Market Has Slower Cumulative Response

Suppose broad market ETFs show smaller short-run coefficients but a more persistent positive response through lags 1–6. This would suggest slower but steadier allocation behavior.

18.3 Case 3: Thematic ETFs Show Convexity

Suppose thematic ETFs display a weak average linear effect but a strong positive quadratic term. This would suggest that only very strong trailing returns generate pronounced inflow response.

18.4 Case 4: Relative Performance Dominates

Suppose absolute returns become less significant after adding category-relative returns. This would imply that investors allocate based more on peer outperformance than on raw returns alone.

19 Potential Deliverables for NASDAQ ETF Intelligence / Indexes

This project could yield several useful outputs.

19.1 Internal Research Note

A polished empirical note on investor demand dynamics across ETF exposures.

19.2 Exposure-Level Flow Elasticity Dashboard

For each exposure category:

- immediate flow elasticity,
- 3-month cumulative elasticity,
- 6-month cumulative elasticity,
- time-to-half-response,
- nonlinear sensitivity.

19.3 Fund-Level Predicted Flow Signal

A model-based score ranking ETFs by expected near-term demand given recent return history and category-specific investor behavior.

19.4 Index/Product Intelligence

Insights into which exposures are more likely to attract assets after strong performance and which require sustained performance before flows materialize.

20 Data Table Blueprint

Variable	Type	Definition	Why Useful
Fund ID	Identifier	Unique ETF identifier	Panel estimation
Date	Time index	Month-end date	Monthly panel structure
Net Flow	Numeric	Dollar creations minus redemptions	Primary dependent input
Lagged AUM	Numeric	Prior month AUM	Flow normalization
Flow Rate	Numeric	Flow / lagged AUM	Main dependent variable
1M Return	Numeric	Monthly total return	Short-horizon signal
3M Return	Numeric	Rolling 3-month total return	Momentum input
6M Return	Numeric	Rolling 6-month total return	Core project signal
12M Return	Numeric	Rolling 12-month total return	Longer-horizon persistence
Category	Categorical	Exposure classification	Cross-exposure comparison
Relative Return	Numeric	Return minus category benchmark/median	Peer outperformance effect
Volatility	Numeric	Rolling realized vol	Risk sensitivity
Drawdown	Numeric	Rolling max drawdown	Loss sensitivity / asymmetry
Log AUM	Numeric	log(AUM)	Scale control
Fund Age	Numeric	Months since launch	Lifecycle effect
Expense Ratio	Numeric	Annual fee	Product attractiveness control
ADV / Liquidity	Numeric	Average daily volume or similar	Tradability / adoption friction
Issuer	Categorical	ETF sponsor	Brand effect
Active/Passive	Binary	Structure indicator	Structural control
Macro Regime	Categorical	Inflation / drawdown / high-vol state	Conditional response analysis

21 Recommended Empirical Design Choices

21.1 Frequency

Monthly data is the best initial choice. It reduces noise, aligns with many allocation decisions, and is easier to interpret than daily flow reactions. Weekly data can be explored later.

21.2 Sample Horizon

Ten years is a reasonable starting window, long enough to capture multiple market regimes:

- low-rate expansion,
- inflationary regime,
- crisis periods,
- commodity surges,
- growth/value rotations.

21.3 Category Construction

Use internally meaningful categories with adequate sample size. Avoid overly broad bins that dilute behavior and avoid overly narrow bins that create unstable estimates.

21.4 Inference

Use clustered standard errors. For robustness, compare:

- equal-weighted regressions,
- AUM-weighted regressions,
- subsamples by size tier.

22 Why This Project Matters Strategically

This project is useful because it links *market performance*, *investor demand*, and *product exposure* in a unified quantitative framework. It goes beyond descriptive flow commentary and moves toward a measurable dynamic demand model.

For NASDAQ ETF Intelligence / Indexes, this can help answer:

- which exposures attract tactical demand fastest,
- which categories require persistent performance to gather assets,
- how investor response differs in defensive vs. risk-on exposures,
- whether certain index-linked segments have latent asset-gathering potential after strong returns.

It can also support communication with issuers, inform product research, and potentially guide internal views on exposure viability and timing sensitivity.

23 Final Recommendation

The project should begin with a clean, high-quality, monthly panel over the past 10 years and proceed in the following order:

1. Build normalized flow measures.
2. Run descriptive bucket charts of future flows against trailing returns by exposure.
3. Estimate fixed-effects models using trailing 1M, 3M, and 6M returns.
4. Upgrade to distributed-lag regressions to estimate response speed.
5. Compare cumulative response curves across exposures.
6. Add nonlinear and relative-performance extensions.
7. Construct a flow elasticity score and predicted-flow momentum signal.

The central empirical output should be a set of exposure-specific dynamic response curves:

$$\{\widehat{\beta}_{e,1}, \widehat{\beta}_{e,2}, \ldots, \widehat{\beta}_{e,K}\},$$

along with cumulative elasticity measures:

$$\widehat{\mathcal{E}}_e(1), \quad \widehat{\mathcal{E}}_e(3), \quad \widehat{\mathcal{E}}_e(6).$$

These quantities translate the initial intuition—that gold may react faster than broad market exposures—into a formal, testable, and economically meaningful research framework.

24 Very Simple Project Roadmap

This section provides a simplified implementation roadmap translating the research framework into a practical sequence of steps. The purpose is to make the project easy to execute in order, while linking each implementation step back to the relevant methodological section of this paper.

24.1 Step 1: Build the ETF dataset

The first step is to create a clean monthly panel dataset of ETFs over roughly the past 10 years. For each fund and month, we should collect:

- net flows,
- AUM,
- returns,
- category / exposure type,
- basic fund characteristics such as age, expense ratio, issuer, and liquidity.

This is the foundation of the entire project. Without a clean fund-level panel, none of the later modeling can be done properly.

Reference in paper:

- Section 6: Data Requirements
- Section 7: Flow Construction and Measurement Issues
- Section 18: Data Table Blueprint

24.2 Step 2: Create the main project variables

Once the raw data is assembled, the next step is to create the variables used in the analysis. The most important are:

- normalized flow rate,
- trailing 1M, 3M, 6M, and 12M returns,
- recent vs. older return blocks,
- category-relative return measures,
- controls such as size, volatility, and fund age.

In simple terms, this is where we convert raw ETF data into the variables needed to test whether flows follow returns.

Reference in paper:

- Section 6.2: Derived Variables
- Section 7: Flow Construction and Measurement Issues

24.3 Step 3: Do simple descriptive analysis first

Before running regressions, we should first look at the data in a simple and intuitive way. For example:

- sort funds into buckets based on trailing 6M return,
- calculate average next-month flow for each bucket,
- compare the patterns across exposure types such as gold and broad market equity.

This helps us see whether the relationship is positive, nonlinear, fast, slow, or different across exposures before introducing heavier econometric structure.

Reference in paper:

- Section 8: Initial Exploratory Framework

24.4 Step 4: Run the basic regression

After the descriptive stage, the next step is to estimate a simple panel regression where future flows are explained by trailing returns and basic controls.

This gives the first formal answer to the question: do stronger past returns predict stronger future ETF flows?

Reference in paper:

- Section 9: Baseline Econometric Specification
- Section 17.1: Specification A
- Section 17.2: Specification B

24.5 Step 5: Run the distributed-lag regression

Once the baseline model is working, the main model should be the distributed-lag regression. This allows us to measure not just whether flows respond to returns, but how quickly they respond.

This is probably the most important section of the project, because it turns the original intuition into a measurable response curve.

Reference in paper:

- Section 10: Distributed-Lag Regression
- Section 17.3: Specification C

24.6 Step 6: Compare categories and exposure types

After estimating the flow-response model, the next step is to compare the coefficients across exposures. For example:

- does gold respond faster than broad market ETFs?
- do thematic funds respond more sharply?
- do fixed income ETFs react differently from equities?

This step is where the project becomes strategically interesting, because it moves from one average ETF effect to a cross-exposure demand map.

Reference in paper:

- Section 11: Exposure-Specific Estimation
- Section 17.4: Specification D

24.7 Step 7: Add relative performance and nonlinear effects

Once the core model is established, we can make the project more sophisticated by testing:

- whether outperformance versus category peers matters more than absolute returns,
- whether very strong returns produce disproportionately large flows,
- whether positive and negative returns have asymmetric effects.

This step is not necessary for the first version, but it is the natural extension once the baseline results are stable.

Reference in paper:

- Section 12: Relative Performance vs. Absolute Performance
- Section 13: Nonlinearity and Threshold Effects
- Section 17.5: Specification E

24.8 Step 8: Build the flow elasticity / momentum score

After the coefficients are estimated, we can convert them into a practical output:

- an exposure-level flow elasticity measure,
- a cumulative response score,
- a predicted-flow or momentum score for funds or categories.

This is the stage where the project becomes directly usable for ETF intelligence, issuer analysis, or internal dashboarding.

Reference in paper:

- Section 13: A Mathematical Definition of Flow Elasticity

24.9 Step 9: Turn the model into an internal research product

The final step is to present the results in a way that is useful for business and research purposes. This could include:

- category rankings by speed of flow response,
- charts of cumulative elasticity by exposure,
- a dashboard showing expected investor demand after strong recent performance.

This makes the project useful not just as an academic exercise, but as a practical ETF intelligence framework.

Reference in paper:

- Section 19: Recommended Empirical Design Choices
- Section 20: Why This Project Matters Strategically
- Section 21: Final Recommendation

24.10 Simple Summary of the Order

If the project is done in the simplest practical order, it should go as follows:

1. pull ETF data,
2. build normalized flow and return variables,
3. do simple return-bucket charts,
4. run baseline regressions,
5. run distributed-lag regressions,
6. compare exposure types,
7. add relative performance and nonlinear extensions,
8. create a momentum / flow elasticity score,
9. package results into internal research outputs.

Appendix: Compact Executive Summary

- The project studies how ETF flows respond to prior returns.
- The key idea is **flow elasticity**: how sensitive and how fast flows react.
- The best framework is a **distributed-lag regression** on a fund-level monthly panel.
- Compare categories such as gold vs. broad market equity.
- Use normalized flows, trailing returns, category controls, and fund characteristics.
- Add relative performance and nonlinear effects as phase 2.
- Convert estimated dynamic response coefficients into a practical **momentum / predicted-flow score**.